



Randomized LU decomposition



Gil Shabat^a, Yaniv Shmueli^b, Yariv Aizenbud^c, Amir Averbuch^{b,*}

^a School of Electrical Engineering, Tel Aviv University, Tel Aviv 69978, Israel

^b School of Computer Science, Tel Aviv University, Tel Aviv 69978, Israel

^c Department of Applied Mathematics, School of Mathematical Sciences, Tel Aviv University, Tel Aviv 69978, Israel

ARTICLE INFO

Article history:

Received 20 February 2015

Received in revised form 17 April 2016

Accepted 29 April 2016

Available online 5 May 2016

Communicated by Thomas Strohmer

Keywords:

LU decomposition
Matrix factorizations
Random matrices
Randomized algorithms

ABSTRACT

Randomized algorithms play a central role in low rank approximations of large matrices. In this paper, the scheme of the randomized SVD is extended to a randomized LU algorithm. Several error bounds are introduced, that are based on recent results from random matrix theory related to subgaussian matrices. The bounds also improve the existing bounds of already known randomized SVD algorithm. The algorithm is fully parallelized and thus can utilize efficiently GPUs without any CPU–GPU data transfer. Numerical examples, which illustrate the performance of the algorithm and compare it to other decomposition methods, are presented.

© 2016 Elsevier Inc. All rights reserved.

1. Introduction

Matrix factorizations and their relations to low rank approximations play a major role in many of today's applications [1]. In mathematics, matrix decompositions are used for low rank matrix approximations that often reveal interesting properties in a matrix. Matrix decompositions are used for example in solving linear equations and in finding least squares solutions. In engineering, matrix decompositions are used in computer vision [2], machine learning [3], collaborative filtering and Big Data analytics [4]. As the size of the data grows exponentially, feasible methods for the analysis of large datasets have gained an increasing interest. Such an analysis can involve a factorization step of the input data given as a large sample-by-feature matrix or given by a sample affinity matrix [5–7]. High memory consumption and the computational complexity of the factorization step are two main reasons for the difficulties in analyzing huge data structures. Recently, there is an on-going interest in applying mathematical tools that are based on randomization to overcome these difficulties.

* Corresponding author. Fax: +972 3 6422020.

E-mail address: amir@math.tau.ac.il (A. Averbuch).

Some of the randomized algorithms use random projections that project the matrix to a set of random vectors. Formally, given a matrix A of size $m \times n$ (assume $m \geq n$) and a random matrix G of size $n \times k$, then the product AG is computed to obtain a smaller matrix that potentially captures most of the range of A . In most of these applications, k is set to be much smaller than n to obtain a compact approximation for A .

In this paper, we develop a randomized version of the LU decomposition. Given an $m \times n$ matrix A , we seek a lower triangular $m \times k$ matrix L and an upper triangular $k \times n$ matrix U such that

$$\|LU - PAQ\|_2 = C(m, n, k)\sigma_{k+1}(A), \quad (1.1)$$

where P and Q are orthogonal permutation matrices, $\sigma_{k+1}(A)$ is the $k+1$ largest singular value of A and $C(m, n, k)$ is a constant that depends on m, n and k .

The interest in a randomized LU decomposition can be motivated (computationally-wise) by two important properties of the classical LU decomposition: First, it can be applied efficiently to sparse matrices with computation time that depends on the number of non-zero elements. LU decomposition with full pivoting on sparse matrices can generate large regions of zeros in the factorized matrices [8–10]. Processing of sparse matrices will be treated in a separate paper. Second, LU decomposition can be fully parallelized [11] which makes it applicable for running on Graphics Processing Units (GPU). GPUs are mostly used for computer games, graphics and visualization such as movies and 3D display. Their powerful computation capabilities can be used for fast matrix computations [12].

The contributions of the paper are twofold: A randomized version for LU decomposition, which is based on the randomized SVD template [13,14], is presented. The algorithm is analyzed and several error bounds are derived. The bounds are based on recent results from random matrix theory for the largest and smallest singular values of random matrices with subgaussian entries [15,16]. This technique is also used to improve the bounds for the randomized SVD. The randomized LU is fully implemented to run on a standard GPU card without any GPU–CPU data transfer. It enables us to accelerate the algorithm significantly. We present numerical results that compare our algorithm with other decomposition methods and show that it outperforms them.

The paper is organized as follows: Section 2, overviews related work on matrix decomposition and approximation that use randomized methods. Section 3 reviews several mathematical results that are needed for the development of the randomized LU. Section 4 presents several randomized LU algorithms and several error bounds on their approximations are proved. Section 5 presents numerical results on the approximation error, the computational complexity of the algorithm and compares it with other methods. The performance comparison was done on different types of matrices and by using GPU cards.

2. Related work

Efficient matrix decomposition serves as a basis for many studies and algorithms design for data analysis and applications. Fast randomized matrix decomposition algorithms are used for tracking objects in videos [7], multiscale extensions for data [17] and detecting anomalies in network traffic for finding cyber attacks [18], to name some. There are randomized versions for many different matrix factorization algorithms [14], compressed sensing [19] and least squares [20].

There is a variety of methods and algorithms that factorize a matrix into several matrices. Typically, the factorized terms have properties such as being triangular, orthogonal, diagonal, sparse or low rank. In general, a certain control on the desired approximation error for a factorized matrix is possible. For example, it is achievable by increasing the rank of a low rank approximation or by allowing dense factors for sparse decompositions.

Rank revealing factorization uses permutation matrices on the columns and rows of a matrix A so that the factorized matrices structure has a strong rank portion and a rank deficient portion. The most known example for approximating an $m \times n$ matrix A by a low rank k matrix is the truncated SVD. Other rank revealing factorizations can be used to achieve low rank approximations. For example, both QR and LU factorizations have rank revealing versions such as RRQR decomposition [21], strong RRQR [22] decomposition, RRLU decomposition [23] and strong RRLU decomposition [24].

Other matrix factorization methods such as Interpolative Decomposition (ID) [25] and CUR decomposition [26], use columns and rows of the original matrix A in the factorization process. Such a property exposes the most important portions that construct A . An ID factorization of order k of an $m \times n$ matrix A consists of an $m \times k$ matrix B whose columns consist of a subset of the columns of A , as well as a $k \times n$ matrix P , such that a subset of the columns of P becomes a $k \times k$ identity matrix and $A \approx BP$ such that $\|A - BP\| \lesssim \mathcal{O}(n, \sigma_{k+1}(A))$. Usually, $k = \#\{j : \sigma_j(A) \geq \delta \sigma_1(A)\}$ is the numerical rank of A up to a certain accuracy $\delta > 0$. This selection of k guarantees that the columns of B constitute a well-conditioned basis for the range of A [25].

Randomized version for many important algorithms has been developed in order to reduce the computational complexity by approximating the solution to a desired rank. These include SVD, QR and ID factorizations [13], CUR decomposition as a randomized version [26] of the pseudo-skeleton decomposition, methods for solving least squares problems [27,28,20] and low rank approximations [28,29].

Randomized algorithms for computing matrix factorizations typically employ to steps: 1. A low-dimensional space, which captures most of the “energy” of A , is found using randomization. 2. A is projected onto the retrieved subspace and the projected matrix is factorized [14].

Several different options exist when random projection matrix is used in the step 1. For example, it can be a matrix of random signs (± 1) [30,31], a matrix of i.i.d. Gaussian random variables with zero mean and unit variance [13], a matrix whose columns are selected randomly from the identity matrix with either uniform or non-uniform probability [32,33], a random sparse matrix designed to enable fast multiplication with a sparse input matrix [28,29], random structured matrices, which use orthogonal transforms such as discrete Fourier transform, Walsh–Hadamard transform and more [27,20,34]. In our algorithm, we use Gaussian matrices in Step 1 as well as structured Fourier matrices to achieve accelerated computation.

3. Preliminaries

In this section, we review the rank revealing LU (RRLU) decomposition and bounds on singular values bounds for random matrices that will be used to prove the error bounds for the randomized LU algorithm. Throughout the paper, we use the following notation: for any matrix A , $\sigma_j(A)$ is the j th largest singular value and $\|A\|$ is the spectral norm (the largest singular value or l_2 operator norm). If x is a vector then $\|x\|$ is the standard l_2 (Euclidean) norm. A^\dagger denotes the pseudo-inverse of A . For a random variable X , \mathbb{E} denotes the expectation of X and $\mathbb{P}(X \geq x)$ is the probably of a random variable X to be larger than a scalar x .

3.1. Rank revealing LU (RRLU)

The following theorem is adapted from [23] (Theorem 1.2):

Theorem 3.1 ([23]). *Let A be an $m \times n$ matrix ($m \geq n$). Given an integer $1 \leq k < n$, the following factorization*

$$PAQ = \begin{pmatrix} L_{11} & 0 \\ L_{21} & I_{n-k} \end{pmatrix} \begin{pmatrix} U_{11} & U_{12} \\ 0 & U_{22} \end{pmatrix} \quad (3.1)$$

holds where L_{11} is a unit lower triangular, U_{11} is an upper triangular, P and Q are orthogonal permutation matrices. Let $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n \geq 0$ be the singular values of A , then

$$\sigma_k \geq \sigma_{\min}(L_{11}U_{11}) \geq \frac{\sigma_k}{k(n-k)+1}, \quad (3.2)$$

and

$$\sigma_{k+1} \leq \|U_{22}\| \leq (k(n-k)+1)\sigma_{k+1}. \quad (3.3)$$

This is called RRLU decomposition. Based on [Theorem 3.1](#), we have the following definition:

Definition 3.1 (*RRLU rank k approximation denoted $RRLU_k$*). Given a RRLU decomposition ([Theorem 3.1](#)) of a matrix A with an integer k (as in Eq. (3.1)) such that $PAQ = LU$. The RRLU rank k approximation is defined by taking k columns from L and k rows from U such that

$$RRLU_k(PAQ) = \begin{pmatrix} L_{11} \\ L_{21} \end{pmatrix} (U_{11}U_{12}) \quad (3.4)$$

where $L_{11}, L_{21}, U_{11}, U_{12}, P$ and Q are defined in [Theorem 3.1](#).

Lemma 3.2 (*RRLU approximation error*). The error of the $RRLU_k$ approximation of A is

$$\|PAQ - RRLU_k(PAQ)\| \leq (k(n-k)+1)\sigma_{k+1}. \quad (3.5)$$

Proof. The proof follows directly from Eqs. (3.1) and (3.4). \square

[Lemma 3.3](#) appears in [\[35\]](#), page 75:

Lemma 3.3 ([\[35\]](#)). Let A and B be two matrices and let $\sigma_j(\cdot)$ denote the j th singular value of a matrix. Then, $\sigma_j(AB) \leq \|A\|\sigma_j(B)$ and $\sigma_j(AB) \leq \|B\|\sigma_j(A)$.

[Lemma 3.4](#) was taken from [\[13\]](#) and it is an equivalent formulation for Eq. (8.8) in [\[36\]](#).

Lemma 3.4 ([\[13\]](#)). Suppose that G is a real $n \times l$ matrix whose entries are i.i.d. Gaussian random variables with zero mean and unit variance and let m be an integer such that $m \geq l$, $m \geq n$, $\gamma > 1$ and

$$1 - \frac{1}{4(\gamma^2 - 1)\sqrt{\pi m \gamma^2}} \left(\frac{2\gamma^2}{e^{\gamma^2 - 1}} \right)^m \geq 0. \quad (3.6)$$

Then, $\|G\| \leq \sqrt{2m\gamma}$ with probability not less than the value in Eq. (3.6).

3.2. Subgaussian random matrices

Definition 3.2. A real valued random variable X is called subgaussian if there exists $b > 0$ such that for all $t > 0$ we have $\mathbb{E}e^{tX} \leq e^{b^2 t^2 / 2}$.

We review several results adapted from [\[15,37\]](#) about random matrices whose entries are subgaussian. We focus on the case where A is a tall $m \times n$ matrix ($m > (1 + \frac{1}{\ln n})n$). Similar results can be found in [\[16\]](#) for square and almost square matrices.

Definition 3.3. Assume that $\mu \geq 1$, $a_1 > 0$ and $a_2 > 0$. $\mathcal{A}(\mu, a_1, a_2, m, n)$ is the set of all $m \times n$ ($m > n$) random matrices $A = (\xi_{ij})$ whose entries are centered i.i.d. real valued random variables satisfying the following conditions:

1. Moments: $\mathbb{E}|\xi_{ij}|^3 \leq \mu^3$;
2. Norm: $\mathbb{P}(\|A\| > a_1\sqrt{m}) \leq e^{-a_2m}$ where \mathbb{P} is a probability function;
3. Variance: $\mathbb{E}\xi_{ij}^2 \geq 1$.

It is shown in [15] that if A is subgaussian then $A \in \mathcal{A}$. For a Gaussian matrix with zero mean and unit variance, $\mu = \left(\frac{4}{\sqrt{2\pi}}\right)^{\frac{1}{3}}$.

Theorems 3.5 and 3.6 are taken from Section 2 in [15].

Theorem 3.5 ([15]). Every matrix A of size $m \times n$ ($m \geq n$), whose entries are subgaussian with $\mu \geq 1$ and $a_2 \geq 0$, satisfies:

$$\mathbb{P}(\|A\| \geq a_1\sqrt{m}) \leq e^{-a_2m} \quad (3.7)$$

where $a_1 = 6\mu\sqrt{a_2 + 4}$.

Theorem 3.5 provides an upper bound for the largest singular value that depends on the desired probability. Theorem 3.6 is used to bound from below the smallest singular value of random Gaussian matrices.

Theorem 3.6 ([15]). Let $\mu \geq 1$, $a_1, a_2 > 0$. Let A be an $m \times n$ matrix where $m > (1 + \frac{1}{\ln n})n$. m can be written as $m = (1 + \delta)n$. Suppose that the entries of A are centered independent random variables such that conditions 1, 2, 3 in Definition 3.3 hold. Then, there exist positive constants c_1 and c_2 such that

$$\mathbb{P}(\sigma_n(A) \leq c_1\sqrt{m}) \leq e^{-m} + e^{-c''m/(2\mu^6)} + e^{-a_2m} \leq e^{-c_2m}. \quad (3.8)$$

From Theorem 3.6, the exact values of constants c_1, c_2 and c'' are

$$c_1 = \frac{b}{e^2 c_3} \left(\frac{b}{3e^2 c_3 a_1} \right)^{\frac{1}{\delta}}, \quad c'' = \frac{27}{2^{11}} \quad (3.9)$$

where $c_3 = 4\sqrt{\frac{2}{\pi}} \left(\frac{2\mu^9}{a_1^3} + \sqrt{\pi} \right)$, $b = \min\left(\frac{1}{4}, \frac{c'}{5a_1\mu^3}\right)$ and $c' = \left(\frac{27}{2^{13}}\right)^{\frac{1}{2}}$. For the constant c_2 , we need a small enough constant to satisfy the inequality in Eq. (3.8) and set it, for simplification, to

$$c_2 = \min\left(1, \frac{c''}{(2\mu^6)}, a_2\right) - \frac{\ln 3}{m}. \quad (3.10)$$

The setting of c_2 according to Eq. (3.10) comes from a relaxation of the inequality

$$e^{-m} + e^{-c''m/(2\mu^6)} + e^{-a_2m} \leq 3e^{-\min\left(1, \frac{c''}{(2\mu^6)}, a_2\right)m} \leq e^{-c_2m}$$

and solving $3e^{-\min\left(1, \frac{c''}{(2\mu^6)}, a_2\right)m} \leq e^{-c_2m}$ for c_2 .

3.3. The SRFT matrix

The Subsampled Random Fourier Transform (SRFT), which is described in [38,39], is a random matrix R with the structure $R = DFS$ where D is an $n \times n$ diagonal matrix whose entries are i.i.d. random variables

drawn from a uniform distribution on the unit circle in \mathbb{C} , F is an $n \times n$ discrete Fourier transform such that $F_{jk} = \frac{1}{\sqrt{n}} e^{-2\pi i(j-1)(k-1)/n}$ and S is an $n \times l$ matrix whose entries are all zeros except for a single randomly placed 1 in each column.

Lemma 3.7 ([39]). *For any $m \times n$ matrix A , let R be the $n \times l$ SRFT matrix. Then, $Y = AR$ can be computed in $\mathcal{O}(mn \log l)$ floating point operations.*

3.4. Interpolative decomposition (ID)

Let A be an $m \times n$ of rank r . $A \approx A_{(:,J)}X$ is the ID of rank r of A if:

1. J is a subset of r indices from $1, \dots, n$.
2. The $r \times n$ matrix $A_{(:,J)}$ is a subset of J columns from A .
3. X is an $r \times n$ matrix whose entries are less than 2 in magnitude and contains r columns of the identity matrix.

Similarly, it is possible to compute the ID with row selection such that $A \approx XA_{(J,:)}$. The ID is based on [22] and it is introduced in [13,25,14] for deterministic and random algorithms. It is possible to compute ID with LU instead of using QR. This can increase the reconstruction error, since RRQR has better bounds than RRLU [22,23] while reducing the computational complexity since LU is faster to compute than QR [11].

4. Randomized LU

In this section, we present the randomized LU algorithm (Algorithm 4.1) that computes the LU rank k approximation of a full matrix. In addition, we present Algorithm 4.4 that utilizes the SRFT matrix for achieving a faster processing. Error bounds are derived for each algorithm.

The algorithm begins by projecting the input matrix on a random matrix. The resulting matrix captures most of the range of the input matrix. Then, we compute a triangular basis for this matrix and project the input matrix on it. Finally, we find a second triangular basis for the projected columns and multiply it with the original basis. The product leads to a lower triangular matrix L and the upper triangular matrix U is obtained from the second LU factorization.

Remark 4.1. The pseudo-inverse of L_y in step 5 can be computed by $L_y^\dagger = (L_y^T L_y)^{-1} L_y^T$. This can be done efficiently when it is computed on platforms such as GPUs that can multiply matrices via parallelization. Usually, the inversion is done on a small matrix since in many cases $k \ll n$ and therefore it can be done cheaply (computationally wise) by the application of Gaussian elimination.

Algorithm 4.1: Randomized LU decomposition.

Input: A matrix of size $m \times n$ to decompose, k desired rank, $l \geq k$ number of columns to use.

Output: Matrices P, Q, L, U such that $\|PAQ - LU\| \leq \mathcal{O}(\sigma_{k+1}(A))$ where P and Q are orthogonal permutation matrices, L and U are the lower and upper triangular matrices, respectively.

- 1: Create a matrix G of size $n \times l$ whose entries are i.i.d. Gaussian random variables with zero mean and unit standard deviation.
 - 2: $Y \leftarrow AG$.
 - 3: Apply RRLU decomposition (Theorem 3.1) to Y such that $PYQ_y = L_y U_y$.
 - 4: Truncate L_y and U_y by choosing the first k columns and the first k rows, respectively, such that $L_y \leftarrow L_y(:, 1:k)$ and $U_y \leftarrow U_y(1:k, :)$.
 - 5: $B \leftarrow L_y^\dagger PA$.
 - 6: Apply LU decomposition to B with column pivoting $BQ = L_b U_b$.
 - 7: $L \leftarrow L_y L_b$.
 - 8: $U \leftarrow U_b$.
-

Table 4.1

Calculated values for the success probability ξ (Eq. (4.2)).
The terms $l - k$, β and γ appear in Eq. (4.2).

| $l - k$ | β | γ | ξ |
|---------|---------|----------|---------------------------|
| 3 | 5 | 5 | $1 - 6.8 \times 10^{-5}$ |
| 5 | 5 | 5 | $1 - 9.0 \times 10^{-8}$ |
| 10 | 5 | 5 | $1 - 5.2 \times 10^{-16}$ |
| 3 | 30 | 5 | $1 - 5.2 \times 10^{-8}$ |
| 5 | 30 | 5 | $1 - 1.9 \times 10^{-12}$ |
| 10 | 30 | 5 | $1 - 1.4 \times 10^{-24}$ |
| 3 | 30 | 10 | $1 - 5.2 \times 10^{-8}$ |
| 5 | 30 | 10 | $1 - 1.9 \times 10^{-12}$ |
| 10 | 30 | 10 | $1 - 1.4 \times 10^{-24}$ |

Remark 4.2. In practice, it is sufficient to perform step 3 in Algorithm 4.1 using standard LU decomposition with partial pivoting instead of applying RRLU. The cases where U grows exponentially are extremely rare – see section 3.4.5 in [11,40].

Theorem 4.3 presents an error bound for Algorithm 4.1:

Theorem 4.3. Let A be a matrix of size $m \times n$. Then, its randomized LU decomposition produced by Algorithm 4.1 with integers k and l ($l \geq k$) satisfies:

$$\|LU - PAQ\| \leq \left(2\sqrt{2nl\beta^2\gamma^2 + 1} + 2\sqrt{2nl}\beta\gamma(k(n-k) + 1)\right) \sigma_{k+1}(A), \quad (4.1)$$

with probability not less than

$$\xi \triangleq 1 - \frac{1}{\sqrt{2\pi(l-k+1)}} \left(\frac{e}{(l-k+1)\beta}\right)^{l-k+1} - \frac{1}{4(\gamma^2-1)\sqrt{\pi n\gamma^2}} \left(\frac{2\gamma^2}{e\gamma^2-1}\right)^n, \quad (4.2)$$

where $\beta > 0$ and $\gamma > 1$.

The proof of Theorem 4.3 is given in Section 4.2. To show that the success probability ξ in Eq. (4.2) is sufficiently high, we present several calculated values of ξ in Table 4.1. We omitted the value of n from Table 4.1 since it does not affect the value of ξ due to the fact that the second term in Eq. (4.2) decays fast.

In Section 5, we show that in practice, Algorithm 4.1 produces comparable results to other well-known randomized factorization methods of low rank matrices such as randomized SVD and randomized ID.

4.1. Computational complexity analysis

To compute the number of floating points operations in Algorithm 4.1, we evaluate the complexity of each step:

Step 1: Generating an $n \times l$ random matrix requires $\mathcal{O}(nl)$ operations.

Step 2: Multiplying A by G to form Y requires $l\mathcal{C}_A$ operations, where \mathcal{C}_A is the complexity of applying A to an $n \times 1$ column vector.

Step 3: Partial pivoting computation of LU for Y requires $\mathcal{O}(ml^2)$ operations.

Step 4: Selecting the first k columns (we do not modify them) requires $\mathcal{O}(1)$ operations.

Step 5: Computing the pseudo-inverse of L_y requires $\mathcal{O}(k^2m + k^3 + k^2m)$ operations and multiplying it by A requires $k\mathcal{C}_{A^T}$ operations. Note that P is a permutation matrix that does not modify the rows of A .

Step 6: Computing the partial pivoting LU for B requires $\mathcal{O}(k^2n)$ operations.

Step 7: Computing L requires $\mathcal{O}(k^2m)$ operations.

Step 8: Computing U requires $\mathcal{O}(1)$ operations.

By summing up the complexities of all the steps above, then [Algorithm 4.1](#) necessitated

$$\mathcal{C}_{RandLU} = l\mathcal{C}_A + k\mathcal{C}_{A^T} + \mathcal{O}(l^2m + k^3 + k^2n) \quad (4.3)$$

operations. Here, we used \mathcal{C}_A (and \mathcal{C}_{A^T}) to denote the complexity from the application of A (and A^T) to a vector, respectively. For a general A , $\mathcal{C}_A = \mathcal{C}_{A^T} = \mathcal{O}(mn)$.

4.2. Bounds for the randomized LU (proof of [Theorem 4.3](#))

In this section, we prove [Theorem 4.3](#) and provide an additional complementary bound. This is done by finding a basis to the smaller matrix AG , which is achieved in practice by using RRLU. The assumptions are that L is numerically stable so its pseudo-inverse can be computed accurately, there exists a matrix U such that LU is a good approximation to AG and there exists a matrix F such that $\|AGF - A\|$ is small. L is always numerically stable since it has a small condition number [\[41\]](#).

[Lemmas 4.4, 4.5 and 4.6](#) are needed for the proof of [Theorem 4.3](#). [Lemma 4.4](#) states that a given basis L can form a basis for the columns in A by bounding the error $\|LL^\dagger A - A\|$.

Lemma 4.4. Assume that A is an $m \times n$ matrix, L is an $m \times k$ matrix with rank k , G is an $n \times l$ matrix, l is an integer ($l \geq k$), U is a $k \times l$ matrix and F is $l \times n$ ($k \leq m$) matrix. Then,

$$\|LL^\dagger A - A\| \leq 2\|AGF - A\| + 2\|F\|\|LU - AG\|. \quad (4.4)$$

Proof. By using the triangular inequality we get

$$\|LL^\dagger A - A\| \leq \|LL^\dagger A - LL^\dagger AGF\| + \|LL^\dagger AGF - AGF\| + \|AGF - A\|. \quad (4.5)$$

Clearly, the first term can also be bounded by

$$\|LL^\dagger A - LL^\dagger AGF\| \leq \|LL^\dagger\|\|A - AGF\| \leq \|A - AGF\|. \quad (4.6)$$

The second term can be bounded by

$$\|LL^\dagger AGF - AGF\| \leq \|F\|\|LL^\dagger AG - AG\|. \quad (4.7)$$

In addition,

$$\|LL^\dagger AG - AG\| \leq \|LL^\dagger AG - LL^\dagger LU\| + \|LL^\dagger LU - LU\| + \|LU - AG\|. \quad (4.8)$$

Since $L^\dagger L = I$, it follows that $\|LL^\dagger LU - LU\| = 0$ and that $\|LL^\dagger AG - LL^\dagger LU\| \leq \|AG - LU\|$. When combined with [Eq. \(4.8\)](#) we obtain:

$$\|LL^\dagger AG - AG\| \leq 2\|LU - AG\|. \quad (4.9)$$

By substituting [Eq. \(4.9\)](#) in [Eq. \(4.7\)](#) we get

$$\|LL^\dagger AGF - AGF\| \leq 2\|F\|\|LU - AG\|. \quad (4.10)$$

By substituting Eqs. (4.6) and (4.10) in Eq. (4.5) we get

$$\|LL^\dagger A - A\| \leq 2\|AGF - A\| + 2\|F\|\|LU - AG\|. \quad \square \quad (4.11)$$

Lemma 4.5 appears in [13]. It uses a lower bound for the smallest singular value of a Gaussian matrix with zero mean and unit variance. This bound appears in [42].

Lemma 4.5 ([13]). Assume that k, l, m and n are positive integers such that $k \leq l$, $l \leq \min(m, n)$. Assume that A is a real $m \times n$ matrix, G is $n \times l$ matrix whose entries are i.i.d. Gaussian random variables of zero mean and unit variance, β and γ are real numbers, such that $\beta > 0$, $\gamma > 1$ and the quantity

$$1 - \frac{1}{\sqrt{2\pi(l-k+1)}} \left(\frac{e}{(l-k+1)\beta} \right)^{l-k+1} - \frac{1}{4(\gamma^2-1)\sqrt{\pi n\gamma^2}} \left(\frac{2\gamma^2}{e^{\gamma^2-1}} \right)^n \quad (4.12)$$

is non-negative. Then, there exists a real $l \times n$ matrix F such that

$$\|AGF - A\| \leq \sqrt{2nl\beta^2\gamma^2 + 1}\sigma_{k+1}(A) \quad (4.13)$$

and

$$\|F\| \leq \sqrt{l}\beta \quad (4.14)$$

with probability not less than the value in Eq. (4.12).

Lemma 4.6 rephrases Lemma 4.5 by utilizing the bounds that appear in Section 3.2. The proof is close to the argumentation that appears in the proof of Lemma 4.5.

Lemma 4.6. Let A be a real $m \times n$ ($m \geq n$) matrix. Let G be a real $n \times l$ matrix whose entries are Gaussian i.i.d. with zero mean and unit variance. Let k and l be integers such that $l < \min(m, n)$ and $l > \left(1 + \frac{1}{\ln k}\right)k$. We define a_1, a_2, c_1 and c_2 as in Theorem 3.6. Then, there exists a real matrix F of size $l \times n$ such that

$$\|AGF - A\| \leq \sqrt{\frac{a_1^2 n}{c_1^2 l} + 1}\sigma_{k+1}(A), \quad (4.15)$$

and

$$\|F\| \leq \frac{1}{c_1 \sqrt{l}} \quad (4.16)$$

with probability not less than $1 - e^{-c_2 l} - e^{-a_2 n}$.

Proof. We begin by the application of SVD to A such that

$$A = U\Sigma V^T, \quad (4.17)$$

where U is orthogonal $m \times m$ matrix, Σ is $m \times n$ diagonal matrix with non-negative entries and V is orthogonal $n \times n$ matrix. Assume that given V^T and G , suppose that

$$V^T G = \begin{pmatrix} H \\ R \end{pmatrix}, \quad (4.18)$$

where H is $k \times l$ matrix and R is $(n - k) \times l$ matrix. Since G is a Gaussian i.i.d. matrix and V is an orthogonal matrix, then $V^T G$ is also a Gaussian i.i.d. matrix. Therefore, H is a Gaussian i.i.d. matrix. Define $F = PV^T$, where P is a matrix of size $l \times n$ such that $P = \begin{pmatrix} H^\dagger & 0 \end{pmatrix}$. Therefore,

$$F = \begin{pmatrix} H^\dagger & 0 \end{pmatrix} V^T. \quad (4.19)$$

By computing $\|F\|$ using [Theorem 3.6](#), we get

$$\|F\| = \|PV^T\| = \|H^\dagger\| = \|H^T(HH^T)^{-1}\| = \frac{1}{\sigma_k(H)} \leq \frac{1}{c_1\sqrt{l}} \quad (4.20)$$

with probability not less than $1 - e^{-c_2 l}$. Now, we can bound $\|AGF - A\|$. By using Eqs. (4.17), (4.18) and (4.19) we get

$$AGF - A = U\Sigma \left(\begin{pmatrix} H \\ R \end{pmatrix} \begin{pmatrix} H^\dagger & 0 \end{pmatrix} - I \right) V^T. \quad (4.21)$$

We define S to be the $k \times k$ upper-left block of Σ . Let T be the $(n - k) \times (n - k)$ lower-right block. Then,

$$\Sigma \left(\begin{pmatrix} H \\ R \end{pmatrix} \begin{pmatrix} H^\dagger & 0 \end{pmatrix} - I \right) = \begin{pmatrix} S & 0 \\ 0 & T \end{pmatrix} \begin{pmatrix} 0 & 0 \\ RH^\dagger & -I \end{pmatrix} = \begin{pmatrix} 0 & 0 \\ TRH^\dagger & -T \end{pmatrix}.$$

The norm of the last term is:

$$\left\| \begin{pmatrix} 0 & 0 \\ TRH^\dagger & -T \end{pmatrix} \right\|^2 \leq \|TRH^\dagger\|^2 + \|T\|^2. \quad (4.22)$$

Therefore, by using Eqs. (4.21), (4.22) and the fact that $\|T\| = \sigma_{k+1}(A)$, we get

$$\|AGF - A\| \leq \sqrt{\|TRH^\dagger\|^2 + \|T\|^2} \leq \sqrt{\|H^\dagger\|^2 \|R\|^2 + 1} \sigma_{k+1}(A). \quad (4.23)$$

We also know that

$$\|R\| \leq \|V^T G\| = \|G\| \leq a_1 \sqrt{n}$$

with probability not less than $1 - e^{-a_2 n}$. Combining Eq. (4.23) with the fact that $\|H^\dagger\| \leq \frac{1}{c_1\sqrt{l}}$ and $\|R\| \leq a_1\sqrt{n}$ gives

$$\|AGF - A\| \leq \sigma_{k+1}(A) \sqrt{\frac{a_1^2 n}{c_1^2 l} + 1}. \quad \square \quad (4.24)$$

Remark 4.7. In contrast to [Lemma 4.5](#) where $\|AGF - A\| = \mathcal{O}(\sqrt{nl})$, [Lemma 4.6](#) provides the bound $\|AGF - A\| = \mathcal{O}(\sqrt{\frac{n}{l}})$ that is tighter for large values of l .

Remark 4.8. The condition $l > (1 + \frac{1}{\ln k})k$ in [Lemma 4.6](#) has to be satisfied to meet the error bounds. However, there are bounds for the case where H is almost square ($l \approx k$) or square ($l = k$) and they are given in [\[16\]](#).

Proof of Theorem 4.3. The error is given by $\|LU - PAQ\|$ where L, U, P and Q are the outputs from [Algorithm 4.1](#) whose inputs are the matrix A and integers k and l . From Steps 7 and 8 in [Algorithm 4.1](#) we have

$$\|LU - PAQ\| = \|L_y L_b U_b - PAQ\| \quad (4.25)$$

where L_y is the $m \times k$ matrix in step 4 in Algorithm 4.1. By using the fact that $BQ = L_b U_b = L_y^\dagger PAQ$, we get

$$\|LU - PAQ\| = \|L_y L_b U_b - PAQ\| = \|L_y L_y^\dagger PAQ - PAQ\|. \quad (4.26)$$

The application of Lemma 4.4 to Eq. (4.26) gives

$$\begin{aligned} \|LU - PAQ\| &= \|L_y L_y^\dagger PAQ - PAQ\| \\ &\leq 2\|PAQ\tilde{G}F - PAQ\| + 2\|F\|\|L_y U_y - PAQ\tilde{G}\| \end{aligned} \quad (4.27)$$

where U_y is the $k \times n$ matrix in step 4 in Algorithm 4.1. This holds for any matrix \tilde{G} . In particular, it holds for a matrix \tilde{G} that satisfies $Q\tilde{G} = GQ_y$ where G is a random Gaussian i.i.d. matrix. After rows and columns permutations, G becomes \tilde{G} . Therefore, the last term in Eq. (4.27) can be reformulated as $\|L_y U_y - PAQ\tilde{G}\| = \|L_y U_y - PAGQ_y\|$ where G is the random matrix in Algorithm 4.1. By applying Lemmas 3.2 and 3.3 to $\|L_y U_y - PAQ\tilde{G}\|$ we get

$$\begin{aligned} \|L_y U_y - PAQ\tilde{G}\| &= \|L_y U_y - PAGQ_y\| \\ &\leq (k(n-k) + 1)\sigma_{k+1}(AG) \\ &\leq (k(n-k) + 1)\|G\|\sigma_{k+1}(A). \end{aligned} \quad (4.28)$$

Lemma 4.5 provides that $\|PAQ\tilde{G}F - PAQ\| \leq \sqrt{2nl\beta^2\gamma^2 + 1}\sigma_{k+1}(A)$ and $\|F\| \leq \sqrt{l}\beta$. By combining Lemmas 4.5 and 3.4 we get

$$\|LU - PAQ\| \leq \left(2\sqrt{2nl\beta^2\gamma^2 + 1} + 2\sqrt{2nl}\beta\gamma(k(n-k) + 1)\right)\sigma_{k+1}(A), \quad (4.29)$$

which completes the proof. \square

Remark 4.9. The error in Theorem 4.3 may appear large, especially for the case where $k \approx \frac{n}{2}$ and n is large. Yet, we performed extensive numerical experiments showing that the actual error is much smaller when using Gaussian elimination with partial pivoting. Note that the error can decrease by increasing k . Numerical illustrations appear in section 5.

We now present an additional error bound that relies on [15]. Asymptotically, this is a tighter bound for large values of n and l since it contains the term $\sqrt{\frac{n}{l}}$, which is smaller than the term \sqrt{nl} in Theorem 4.3. See also Remark 4.7.

Theorem 4.10. Given a matrix A of size $m \times n$, integers k and l such that $l > \left(1 + \frac{1}{\ln k}\right)k$ and $a_2 > 0$. By the application of Algorithm 4.1 with A, k and l as its input parameters, the randomized LU decomposition satisfies

$$\|LU - PAQ\| \leq \left(2\sqrt{\frac{a_1^2 n}{c_1^2 l} + 1} + \frac{2a_1 \sqrt{n}}{c_1 \sqrt{l}}(k(n-k) + 1)\right)\sigma_{k+1}(A), \quad (4.30)$$

with probability not less than $1 - e^{-a_2 n} - e^{-c_2 l}$. The value of c_1 is given in Eq. (3.9), the value of c_2 is given in Eq. (3.10) and a_1 is given by Theorem 3.5. a_1, c_1 and c_2 depend on a_2 .

Algorithm 4.2: Randomized algorithm with orthogonal basis.**Input:** A matrix of size $m \times n$ to decompose, matrix rank k , $l \geq k$ number of columns to use.**Output:** Matrix Q of size $m \times k$ such that $\|QQ^*A - A\|$ is bounded, and $Q^*Q = I$ are orthogonal permutation matrices, L and U are lower and upper triangular matrices, respectively.1: Create a matrix G of size $n \times l$ whose entries are i.i.d. Gaussian random variables with zero mean and unit standard deviation.2: $Y \leftarrow AG$.3: Construct a matrix U whose columns form an orthonormal basis for the range of Y using SVD.4: Construct a matrix Q by grouping the first k vectors from U .**Proof.** By using steps 5, 6, 7 and 8 in Algorithm 4.1, we get that

$$\|LU - PAQ\| = \|L_y L_y^\dagger PAQ - PAQ\|. \quad (4.31)$$

Then, from Lemma 4.4

$$\|L_y L_y^\dagger PAQ - PAQ\| \leq 2\|PAQ\tilde{G}F - PAQ\| + 2\|F\|\|L_y U_y - PAQ\tilde{G}\|. \quad (4.32)$$

From Lemma 4.6 we get that

$$\|PAQ\tilde{G}F - PAQ\| \leq \sqrt{\frac{a_1^2 n}{c_1^2 l}} + 1\sigma_{k+1}(A). \quad (4.33)$$

By using the same argumentation given in Theorem 4.3, we get

$$\|L_y U_y - PAQ\tilde{G}\| = \|L_y U_y - PAGQ_y\| \leq (k(n-k) + 1)\|G\|\sigma_{k+1}(A) \quad (4.34)$$

where G is the matrix used in Algorithm 4.1 Step 1. By combining Eqs. (4.32), (4.33) and (4.34), and since $\|F\| \leq \frac{1}{c_1 \sqrt{l}}$ and $\|G\| \leq a_1 \sqrt{n}$ (see Lemma 4.6 and Theorem 3.5, respectively), we get that

$$\|LU - PAQ\| \leq 2\sqrt{\frac{a_1^2 n}{c_1^2 l}} + 1\sigma_{k+1}(A) + \frac{2a_1 \sqrt{n}}{c_1 \sqrt{l}} (k(n-k) + 1)\sigma_{k+1}(A). \quad \square \quad (4.35)$$

4.3. Analysis of the bound in Theorem 4.10

In this section, we analyze the bound in Eq. (4.35). Although Eq. (4.35) bounds the randomized LU decomposition error, this bound can be modified to be used for other randomized algorithms. Randomized algorithms use a range approximation step that generates a smaller matrix than the original matrix that approximates the range of the original matrix. The range approximation enables us to compute the approximated decomposition using a smaller matrix while maintaining a bounded error. The obtained error of a randomized algorithm depends on the quality of the range approximation step. Formally, range approximation of a matrix A can be accomplished by finding an orthogonal matrix Q such that $\|QQ^*A - A\|$ is bounded. Hence, Q^*A is a smaller matrix than A that approximates the range of A . A randomized algorithm, which finds such an orthogonal basis, appears in [14] and described in Algorithm 4.2. Steps 1–3 of the randomized LU decomposition (Algorithm 4.1) employ similar approach, except the basis found is not orthogonal.

By estimating $\|QQ^*A - A\|$ in the same way as was done in Eq. (4.35) and by using Lemma 4.4 instead of estimating $\|LL^\dagger A - A\|$ we get

$$\|QQ^*A - A\| \leq 2\sqrt{\frac{a_1^2 n}{c_1^2 l}} + 1\sigma_{k+1}(A) + \frac{2a_1 \sqrt{n}}{c_1 \sqrt{l}} \sigma_{k+1}(A) \quad (4.36)$$

with probability not less than $1 - e^{-a_2 n} - e^{-c_2 l}$. The value of c_1 is given in Eq. (3.9), the value of c_2 is given in Eq. (3.10) and the value of a_2 is given in Theorem 3.5. All of them depend on a_2 .

Equation (4.36) provides an alternative bound to the randomized SVD algorithm. By neglecting constants and by analyzing the asymptotic behavior of Eq. (4.36) we get that for $n \gg l$

$$\|QQ^*A - A\| \leq 2\sqrt{\frac{a_1^2 n}{c_1^2 l}} + 1\sigma_{k+1}(A) + \frac{2a_1\sqrt{n}}{c_1\sqrt{l}}\sigma_{k+1}(A) \propto \sqrt{\frac{n}{l}}\sigma_{k+1}(A) \quad (4.37)$$

with an asymptotic failure probability of $e^{-c_2 l}$. Two bounds are given in [14]: 1. Expectation-based bound and 2. Probability-based bound. An expectation-based bound that is sharp appears in [43] and the probability bound is better than previously developed bounds in [13]. This probability-based bound is given in [14], Corollary 10.9:

Corollary 4.11 ([14]). For Q from Algorithm 4.2 and $p \geq 4$ ($p = l - k$),

$$\|QQ^*A - A\| \leq (1 + 17\sqrt{1 + k/p})\sigma_{k+1} + \frac{8\sqrt{k+p}}{p+1} \left(\sum_{j>k} \sigma_j^2 \right)^{1/2} \quad (4.38)$$

with failure probability of at most $6e^{-p}$.

We now compare the asymptotic behavior of Eqs. (4.38) with (4.36) for the case of a fixed σ_j , $j > k$, $\sigma_{k+1} = \sigma_{k+2} = \dots = \sigma_{\min(m,n)}$. The asymptotic behavior for $n \gg k + p$ of Eq. (4.38) is given by

$$\|QQ^*A - A\| \leq (1 + 17\sqrt{1 + k/p})\sigma_{k+1} + \frac{8\sqrt{k+p}}{p+1} \left(\sum_{j>k} \sigma_j^2 \right)^{1/2} \propto \frac{\sqrt{(k+p)(n-k)}}{p+1} \sigma_{k+1}. \quad (4.39)$$

Comparison between Eqs. (4.39) and (4.37) shows that Eq. (4.37) provides a better bound since Eq. (4.39) has an additional factor of $\sqrt{k+p}$ in the numerator in comparison to Eq. (4.37) and a smaller denominator than the one in Eq. (4.37). Also, the failure probability is smaller in Eq. (4.37) since the exponents depend on l instead of p .

The bound in Eq. (4.36) is useful especially for large values of l . We assume that $n \gg l$ and $\sigma_j = \sigma$ for $j > k$. Next, we show a numerical example that illustrates the bounds, $m = 2 \cdot 10^8$, $n = 10^8$, $k = 990$, $l = 1000$ and $a_2 = 1$. Computation of a_1 , c_1 and c_2 , which uses Theorems 3.5 and 3.6 provides $a_1 = 15.68$, $c_1 = 0.022$, $c_2 = 0.011$. Substituting these values in Eq. (4.36), provides

$$\|QQ^*A - A\| \leq 2.9 \cdot 10^5 \sigma_{k+1} \quad (4.40)$$

with failure probability $1.1 \cdot 10^{-49}$. The same setup for Eq. (4.38) gives

$$\|QQ^*A - A\| \leq 7.28 \cdot 10^5 \sigma_{k+1} \quad (4.41)$$

with failure probability $2.72 \cdot 10^{-4}$. Clearly, in this example, Eq. (4.36) provides a better bound for both accuracy and failure probability.

Fig. 4.1 compares the asymptotic behaviors of the bounds in Eqs. (4.39) and (4.37). This figure shows that when there is a small oversampling (small p), then the bound in Eq. (4.36), which is indicated by the red line, is asymptotically better in comparison to the bound in Eq. (4.39) which is indicated by the dashed blue line. As the oversampling increases, the bounds coincide. Fig. 4.2 shows the asymptotic behavior of Eq. (4.36) for different values of k and a fixed p . The red line illustrates Eqs. (4.37) and (4.36) and the blue dashed line illustrates Eqs. (4.39) and (4.38).




Fig. 4.1. Bound values vs. oversampling for $k = 3$, $p = 4, 5, \dots, 100$, $l = k + p$.

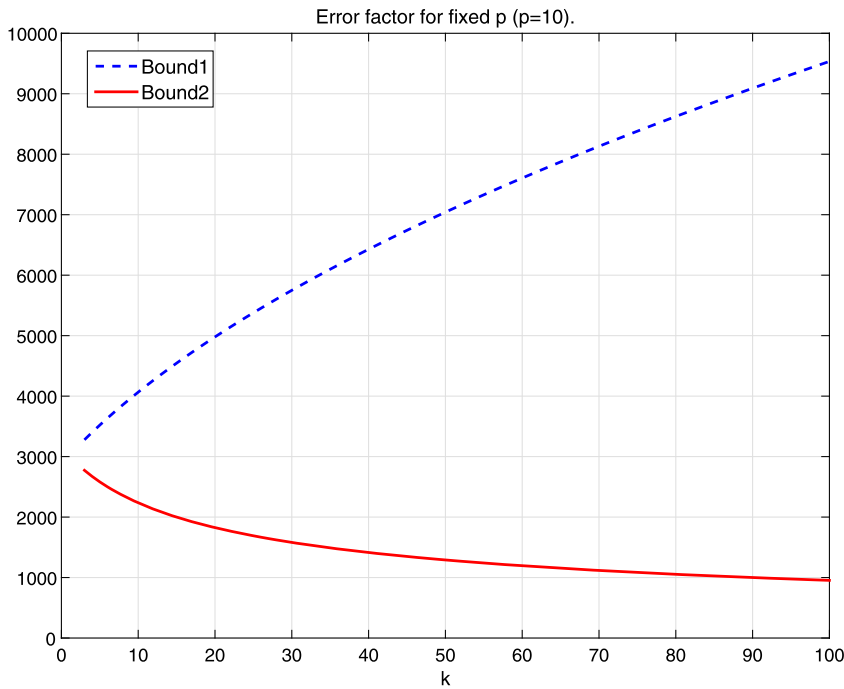


Fig. 4.2. Bound values for fixed $p = 10$ where $k = 3, 4, \dots, 100$.

4.4. Rank deficient least squares

In this section, we use the randomized LU to solve efficiently the Rank Deficient Least Squares (RDLS) problem. Assume that A is an $m \times n$ matrix ($m \geq n$) with $\text{rank}(A) = k$, $k < n$ and b is a column vector of size $m \times 1$. We want to minimize $\|Ax - b\|$. Because A is a rank deficient matrix, then the problem has an infinite number of solutions. We show that the complexity of the solution depends on the rank of A and

Algorithm 4.3: Solving rank deficient least squares with randomized LU.**Input:** Matrix A of size $m \times n$ with rank k , $l, l \geq k$, b vector of size $m \times 1$.**Output:** Solution x that minimizes $\|Ax - b\|$.

- 1: Apply Algorithm 4.1 to A with parameters k and l .
- 2: $y \leftarrow L^\dagger Pb$.
- 3: $z_1 \leftarrow U_1^{-1}y$.
- 4: $z \leftarrow \begin{pmatrix} z_1 \\ z_2 \end{pmatrix}$, where z_2 is an $n - k$ zero vector.
- 5: $x \leftarrow Qz$.

that the problem is equivalent to solving the following two problems: a full rank Least Square (LS) problem of size $m \times k$ and a simplified undetermined linear system of equations that requires a matrix inversion of size $k \times k$. The solution is derived by the application of Algorithm 4.1 to A to get

$$\|Ax - b\| = \|P^T LUQ^T x - b\| = \|LUQ^T x - Pb\|, \quad (4.42)$$

where L is an $m \times k$ matrix, U is a $k \times n$ matrix and both L and U are of rank k . Let $y = UQ^T x$ and $c = Pb$. Then, the problem is reformulated as $\min \|Ly - c\|$. Note that L is a full rank matrix and the problem to be solved becomes a standard full rank LS problem. The solution is given by $y = L^\dagger c$. Next, we solve

$$Uz = y, \quad (4.43)$$

where $z = Q^T x$. Since U is a $k \times n$ matrix, Eq. (4.43) is an underdetermined system. Assume that $U = [U_1 \ U_2]$ and $z = [z_1 \ z_2]^T$, where U_1 is a $k \times k$ matrix, z_1 is a $k \times 1$ vector and z_2 is a $(n - k) \times 1$ vector. Then, the solution is given by setting any value to z_2 and solving $U_1 z_1 = y - U_2 z_2$. For simplicity, we choose $z_2 = 0$. Therefore, we get $z_1 = U_1^{-1}y$. The final solution is given by $x = Qz$. This procedure is summarized in Algorithm 4.3 that finds the solution to the deficient least squares problem that uses Algorithm 4.1.

The complexity of Algorithm 4.3 is equal to the randomized LU complexity (Algorithm 4.1) with an additional inversion cost of the matrix U_1 in Step 3, which is of size $k \times k$. Note that the solution given by Algorithm 4.3 is sparse in the sense that x contains at most k non-zero entries.

4.5. Fast randomized LU

Algorithm 4.1 describes the randomized LU algorithm. This algorithm computes the LU approximation of the matrix A of rank k whose computational complexity is $\mathcal{C}_{RandLU} = \mathcal{O}(lmn + l^2m + k^3 + k^2n)$ operations. We present now an asymptotic improvement to Algorithm 4.1 called fast randomized LU whose computational complexity is

$$\mathcal{C}_{FastRandLU} = \mathcal{O}(mn \log l + mkl + nkl + mk^2 + k^3). \quad (4.44)$$

In order to achieve it, we use the SRFT matrix and the ID Algorithm [25], which were presented in sections 3.3 and 3.4, respectively.

The most computationally expensive procedures are steps 2 and 5 in Algorithm 4.1. Step 2 involves matrix multiplication with the matrix A where A applied to a random matrix. Instead of projecting it with a Gaussian random matrix, we use the SRFT matrix R . Due to the special structure $R = DFS$ (Section 3.3), as was shown in Lemma 3.7, the application of an $m \times n$ matrix A to an $n \times l$ matrix R necessitates $\mathcal{O}(nm \log l)$ floating point operations.

Instead of direct computation of $L_y^\dagger PA$ in step 5 in Algorithm 4.1, A is approximated by the ID of Y , namely, if $Y = XY_{(J,:)}$ is the full rank ID of Y , then $A \approx XA_{(J,:)}$.

Algorithm 4.4: Fast randomized LU decomposition.**Input:** Matrix A of size $m \times n$ to decompose, k desired rank, l number of columns to use.**Output:** Matrices P, Q, L, U such that $\|PAQ - LU\| \leq \mathcal{O}(\sigma_{k+1}(A))$ where P and Q are orthogonal permutation matrices, L and U are the lower and upper triangular matrices, respectively.

- 1: Create a random SRFT matrix R of size $n \times l$ (Lemma 3.7).
- 2: $Y \leftarrow AR$.
- 3: Apply RRLU decomposition to Y such that $PYQ_y = L_yU_y$.
- 4: Truncate L_y and U_y by choosing the first k columns and the first k rows, respectively, such that $L_y \leftarrow L_y(:, 1:k)$ and $U_y \leftarrow U_y(1:k, :)$.
- 5: Compute the full rank ID decomposition of Y such that $Y = XY_{(J,:)}$ (Section 3.4).
- 6: $B \leftarrow L_y^\dagger P X A_{(J,:)}$.
- 7: Apply the LU decomposition to B with column pivoting $BQ = L_bU_b$.
- 8: $L \leftarrow L_yL_b$.
- 9: $U \leftarrow U_b$.

4.5.1. Computational complexity

To compute the number of floating points operations in Algorithm 4.4, we evaluate the complexity of each step:

- Step 1:** The multiplication of an $m \times n$ matrix A by an $n \times l$ matrix R requires $\mathcal{O}(nm \log l)$ operations;
Step 2: The computation of the RRLU decomposition of an $m \times l$ Y requires $\mathcal{O}(ml^2)$ operations;
Step 3: The truncation of L_y and U_y requires $\mathcal{O}(1)$ operations;
Step 4: The computation of the ID decomposition of Y requires $\mathcal{O}(ml^2)$ operations;
Step 5: The computation of the pseudo inverse L_y^\dagger requires $\mathcal{O}(k^2m + k^3)$ operations;
Step 6: The multiplication of $L_y^\dagger P X A_{(J,:)}$ requires $\mathcal{O}(mkl + nkl)$ operations;
Step 7: The computation of the partial pivoting LU of matrix B requires $\mathcal{O}(nk^2)$ operations;
Step 8: The computation of $m \times k$ matrix L requires $\mathcal{O}(mk^2)$ operations.

The total computational complexity of Algorithm 4.4 is $\mathcal{O}(mn \log l + mkl + nkl + mk^2 + k^3)$. By simplifying this expression while assuming that k and l are of the same magnitude, we get that the total computational complexity of Algorithm 4.4 is $\mathcal{O}(mn \log k + (m + n)k^2 + k^3)$.

4.5.2. Correctness Algorithm 4.4

We now prove that Algorithm 4.4 approximates the LU decomposition and provide an error bound.

Theorem 4.12. *Given a matrix A of size $m \times n$. Its fast randomized LU decomposition in Algorithm 4.4 with integers k and l (where $n, m \geq l \geq k$ sufficiently large) satisfies*

$$\|LU - PAQ\| \leq \left(\left[1 + \sqrt{1 + 4k(n - k)} \right] \sqrt{1 + 7n/l} \right) \sigma_{k+1}(A) + 2 \left(\sqrt{\alpha n + 1} + \sqrt{\frac{\alpha}{l}}(k(n - k) + 1) \right) \sigma_{k+1}(A)$$

with probability not less than $1 - 3\frac{1}{\beta k}$ where $\beta > 1$ is a constant.

The proof of Theorem 4.12 uses Lemmas 4.13–4.15.

Lemma 4.13. *Let A be an $m \times n$ matrix with singular values $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_{\min(m,n)}$. Let k and l be integers such that $4 \left[\sqrt{k} + \sqrt{8 \ln(kn)} \right]^2 \ln k \leq l \leq n$. Let R be an $n \times l$ SRFT matrix and $Y = AR$. Denote by Q the $m \times l$ matrix whose columns form an orthonormal basis for the range of Y . Then, with a failure probability of at most $3k^{-1}$ we have*

$$\|A - QQ^*A\| \leq \sqrt{1 + 7n/l} \sigma_{k+1}.$$

[Lemma 4.13](#) appears in [\[14\]](#) as Theorem 11.2 and in [\[44\]](#) as Theorem 3.1 in slightly different formulation.

Lemma 4.14. *Let A be an $m \times n$ matrix, R is an $n \times l$ SRFT random matrix and $Y = XY_{(J,:)}$ is the full rank ID of $Y = AR$. Then,*

$$\|A - XA_{(J,:)}\| \leq \left(1 + \sqrt{1 + 4k(n-k)}\right) \sqrt{1 + 7n/l} \sigma_{k+1}(A)$$

with failure probability of at most $3k^{-1}$ when $l \geq 4 \left(\sqrt{k} + \sqrt{8 \ln(kn)}\right)^2 \ln k$.

The proof is the same as in Lemma 5.1 in [\[14\]](#).

Proof. Denote by Q the matrix whose columns form an orthonormal basis for the range of Y . By using [Lemma 4.13](#) we have

$$\|A - QQ^*A\| \leq \sqrt{1 + 7n/l} \sigma_{k+1}(A) \quad (4.45)$$

except with probability $3k^{-1}$.

Denote $\hat{A} = QQ^*A$. Since $\hat{A} = XQ_{(J,:)}Q^*A$ and $X_{(J,:)} = I$, we have $\hat{A}_{(J,:)} = Q_{(J,:)}Q^*A$. Thus, $\hat{A} = X\hat{A}_{(J,:)}$.

$$\begin{aligned} \|A - XA_{(J,:)}\| &= \|A - X\hat{A}_{(J,:)} + X\hat{A}_{(J,:)} - XA_{(J,:)}\| \\ &\leq \|A - \hat{A}\| + \|X\hat{A}_{(J,:)} - XA_{(J,:)}\| \\ &= \|A - \hat{A}\| + \|X\| \|\hat{A}_{(J,:)} - A_{(J,:)}\| \\ &\leq (1 + \|X\|) \|A - \hat{A}\|. \end{aligned}$$

By using Eq. [\(4.45\)](#) we have

$$\|A - XA_{(J,:)}\| \leq (1 + \|X\|) \sqrt{1 + 7n/l} \sigma_{k+1}(A).$$

The proof is completed since X contains a $k \times k$ identity matrix and the spectral norm of the remaining $(n-k) \times k$ submatrix is bounded by 2. \square

Lemma 4.15 (Appears in [\[39\]](#) as Lemma 4.6). *Suppose that k, l, n and m are positive integers with $k \leq l$ such that $l < \min(m, n)$. Suppose that α and β are real numbers greater than 1 such that*

$$m > l \geq \frac{\alpha^2 \beta}{(\alpha - 1)^2} k^2.$$

Suppose that A is an $m \times n$ complex matrix and Q is the $n \times l$ SRFT matrix. Then, there exists an $l \times n$ complex matrix F such that $\|AQF - A\| \leq \sqrt{\alpha n + 1} \sigma_{k+1}$ and $\|F\| \leq \sqrt{\frac{\alpha}{l}}$ with probability at least $1 - \frac{1}{\beta}$ where σ_{k+1} is the $(k+1)$ th greatest singular value of A .

Lemma 4.16. *Let A, P, Q, L_y and L_y^\dagger be as in [Algorithm 4.4](#), then*

$$\|L_y L_y^\dagger P A Q - P A Q\| \leq 2 \left(\sqrt{\alpha n + 1} + \sqrt{\frac{\alpha}{l}} (k(n-k) + 1) \right) \sigma_{k+1}(A)$$

with probability of at least $1 - \frac{1}{\beta}$ where $m > l \geq \frac{\alpha^2 \beta}{(\alpha - 1)^2} k^2$.

Proof. By applying [Lemma 4.4](#) we get

$$\|L_y L_y^\dagger PAQ - PAQ\| \leq 2\|PAQ\tilde{G}F - PAQ\| + 2\|F\|\|L_y U_y - PAQ\tilde{G}\|. \quad (4.46)$$

U_y is the $k \times n$ matrix in step 4 in [Algorithm 4.4](#). This holds for any matrix \tilde{G} . In particular, for a matrix \tilde{G} satisfies $Q\tilde{G} = RQ_y$, where R is the SRFT matrix in [Algorithm 4.4](#). Therefore, the last term in Eq. (4.46) can be reformulated by $\|L_y U_y - PAQ\tilde{G}\| = \|L_y U_y - PARQ_y\|$. By applying [Lemmas 3.2 and 3.3](#) to $\|L_y U_y - PAQ\tilde{G}\|$ we get

$$\begin{aligned} \|L_y U_y - PAQ\tilde{G}\| &= \|L_y U_y - PARQ_y\| \\ &\leq (k(n-k) + 1)\sigma_{k+1}(AR) \\ &\leq (k(n-k) + 1)\|R\|\sigma_{k+1}(A). \end{aligned} \quad (4.47)$$

Since R is an SRFT matrix, it is orthogonal, thus $\|R\| = 1$. [Lemma 4.15](#) proves that $\|PAQ\tilde{G}F - PAQ\| \leq \sqrt{\alpha n + 1}\sigma_{k+1}(A)$ and $\|F\| \leq \sqrt{\frac{\alpha}{l}}$. By summing up, we get

$$\begin{aligned} \|L_y L_y^\dagger PAQ - PAQ\| &\leq 2\|PAQ\tilde{G}F - PAQ\| + 2\|F\|\|L_y U_y - PAQ\tilde{G}\| \\ &\leq 2\sqrt{\alpha n + 1}\sigma_{k+1}(A) + 2\|F\|(k(n-k) + 1)\sigma_{k+1}(A) \\ &\leq 2(\sqrt{\alpha n + 1} + \sqrt{\frac{\alpha}{l}}(k(n-k) + 1))\sigma_{k+1}(A). \quad \square \end{aligned}$$

Proof of [Theorem 4.12](#).

Proof. By substituting L and U from [Algorithm 4.4](#) we have

$$\begin{aligned} \|LU - PAQ\| &= \|L_y L_b U_b - PAQ\| = \|L_y BQ - PAQ\| = \\ &= \|L_y L_y^\dagger P X A_{(J,:)} Q - PAQ\| \leq \\ &\leq \|L_y L_y^\dagger P X A_{(J,:)} Q - L_y L_y^\dagger PAQ\| + \|L_y L_y^\dagger PAQ - PAQ\|. \end{aligned} \quad (4.48)$$

The first term in the last inequality in Eq. (4.48) is bounded in the following way:

$$\|L_y L_y^\dagger P X A_{(J,:)} Q - L_y L_y^\dagger PAQ\| \leq \|L_y L_y^\dagger P\| \|X A_{(J,:)} - A\| \|Q\| = \|X A_{(J,:)} - A\|.$$

By using [Lemma 4.14](#) we get

$$\|LU - PAQ\| = \|L_y L_b U_b - PAQ\| \leq \left(1 + \sqrt{1 + 4k(n-k)}\right) \sqrt{1 + 7n/l} \sigma_{k+1}(A)$$

with probability of not less than $1 - 3k^{-1}$.

The second term $\|L_y L_y^\dagger PAQ - PAQ\|$ in the last inequality of Eq. (4.48) is bounded by [Lemma 4.16](#).

By combining these results we get

$$\begin{aligned} \|LU - PAQ\| &\leq \left(\left[1 + \sqrt{1 + 4k(n-k)}\right] \sqrt{1 + 7n/l} \right) \sigma_{k+1}(A) \\ &\quad + 2(\sqrt{\alpha n + 1} + \sqrt{\frac{\alpha}{l}}(k(n-k) + 1)) \sigma_{k+1}(A) \end{aligned}$$

which completes the proof. \square




Fig. 5.1. Low rank approximation error of different algorithms: Randomized SVD, randomized ID (QR and LU) and randomized LU with respect to the real singular values of the testing matrix.

5. Numerical results

In order to evaluate [Algorithm 4.1](#), we present the numerical results by comparing the performances of several randomized low rank approximation algorithms. We tested the algorithms and compared them by applying them to random matrices and to images. All the results were computed using the standard MATLAB libraries including MATLAB's GPU interface on a machine with two Intel Xeon CPUs X5560 2.8 GHz that contains an nVidia GPU GTX TITAN card.

5.1. Error rate and computational time comparisons

The performance of the randomized LU ([Algorithm 4.1](#)) was tested and compared to the randomized SVD and to the randomized ID (see [\[13,14\]](#)). The tests compare the normalized (relative) error of the low rank approximation obtained by the examined methods. In addition, the computational time of each method was measured. If A is the original matrix and \hat{A} is a low rank approximation of A , then the relative approximation error is given by:

$$\text{err} = \frac{\|A - \hat{A}\|}{\|A\|}. \quad (5.1)$$

We compared the low rank approximation achieved by the application of the randomized SVD, randomized ID and randomized LU with different ranks k . Throughout the experiments, we chose $l = k + 3$ and the test matrix was a 3000×3000 dense random matrix taken from a normal distribution with zero mean and unit variance. The singular values of the matrix were modified to be exponentially decaying. The computations of the algorithms were done in a single precision. The comparison results are presented in [Fig. 5.1](#). The experiment shows that the error of the randomized ID is larger than the error obtained from both the randomized SVD and the randomized LU ([Algorithm 4.1](#)), which are almost identical. In addition, we compared the execution time of these algorithms. The results are presented in [Fig. 5.2](#). The randomized




Fig. 5.2. The execution times of the same algorithms as in Fig. 5.1 running on a CPU.

LU is fast in comparison due to the fact that QR and SVD factorizations used in the randomized SVD and randomized ID are slower than LU factorization. The LU factorization has a parallel implementation (see [11] section 3.6). To see the impact of the parallel LU decomposition implementation, the execution time to compute the randomized LU of a matrix of size 3000×3000 was measured on an nVidia GTX TITAN GPU device and it is shown in Fig. 5.3. The execution time on the GPU was 10 times faster than running it on an eight cores CPU. Thus, the algorithm scales well. For larger matrices (n and k are large), the differences between the performances while running on CPU and on GPU are more significant.

5.2. Power iterations

The performance (error wise) of Algorithm 4.1 can be further improved by the application of power iterations [14,45] to the input matrix. Specifically, by replacing the projection step $Y \leftarrow AG$ with $Y \leftarrow (AA^*)^q AG$ for small integer q (for larger values, a normalized scheme has to be used – see [45]). Applying the power iterations scheme to an exponentially decaying singular values, an improvement is achieved even for $q = 1$. The same random matrix, which was described in section 5.1, is tested again with and without power iterations on both GPU and CPU. This time, double precision is used. The results are shown in Figs. 5.4 and 5.5.

The experiment was repeated for a slowly decaying singular values. The decay of the singular values was proportional to $1/k^2$. To make the decay slower than $1/k^2$ from $k = 1$, a factor was added such that $\sigma_k = \frac{100}{(9+k)^2}$. The singular values were normalized such that $\sigma_1 = 1$. The results are shown in Fig. 5.6.

5.3. Image matrix factorization

Algorithm 4.1 was applied to images given in a matrix form. The factorization error and the execution time were compared to the performances of the randomized SVD and to the randomized ID. We also added




Fig. 5.3. The execution times from running Algorithm 4.1 on different computational platforms: CPU with 8 cores and GPU.

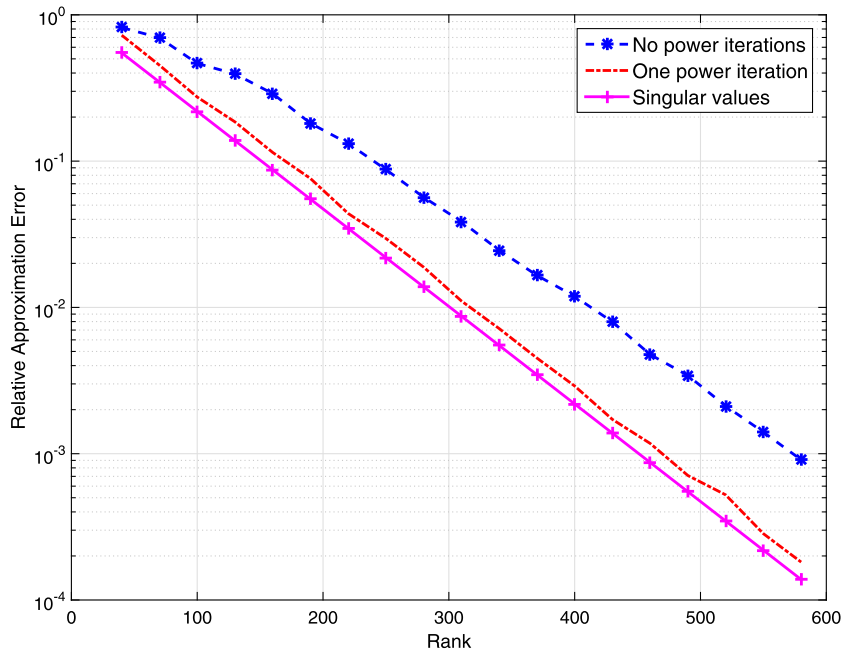


Fig. 5.4. Low rank approximation using randomized LU decomposition (Algorithm 4.1) with and without power iterations.

the SVD error and execution time computed by the Lanczos bidiagonalization [11] that is implemented in the PROPACK package [46]. The image size was 2124×7225 pixels and it has 256 gray levels. The parameters were $k = 200$ and $l = 203$. The approximation quality (error) was measured in PSNR defined by

$$\text{PSNR} = 20 \log_{10} \frac{\max_A \sqrt{N}}{\|A - \hat{A}\|_F} \quad (5.2)$$




Fig. 5.5. The execution time on CPU and GPU of the randomized LU with and without power iterations.




Fig. 5.6. Low rank approximation using randomized LU decomposition with and without power iterations for slowly decaying singular values.

where A is the original image, \hat{A} is the reconstructed (approximated) image, which is the output from Algorithm 4.1, defined by $\hat{A} = P^T L U Q^T$, \max_A is the maximal pixel value of A , N is the total number of pixels and $\|\cdot\|_F$ is the Frobenius norm.

Figs. 5.7 and 5.8 show the original and the reconstructed images, respectively. The image reconstruction quality (measured in PSNR) related to rank k is shown in Fig. 5.9 where for the same k , the PSNR from the application of Algorithm 4.1 is higher than the PSNR generated by the application of the randomized ID




Fig. 5.7. The original input image of size 2124×7225 that was factorized by the application of the randomized LU, randomized ID and randomized SVD algorithms.




Fig. 5.8. The reconstructed image from the application of the randomized LU factorization with $k = 200$ and $l = 203$.

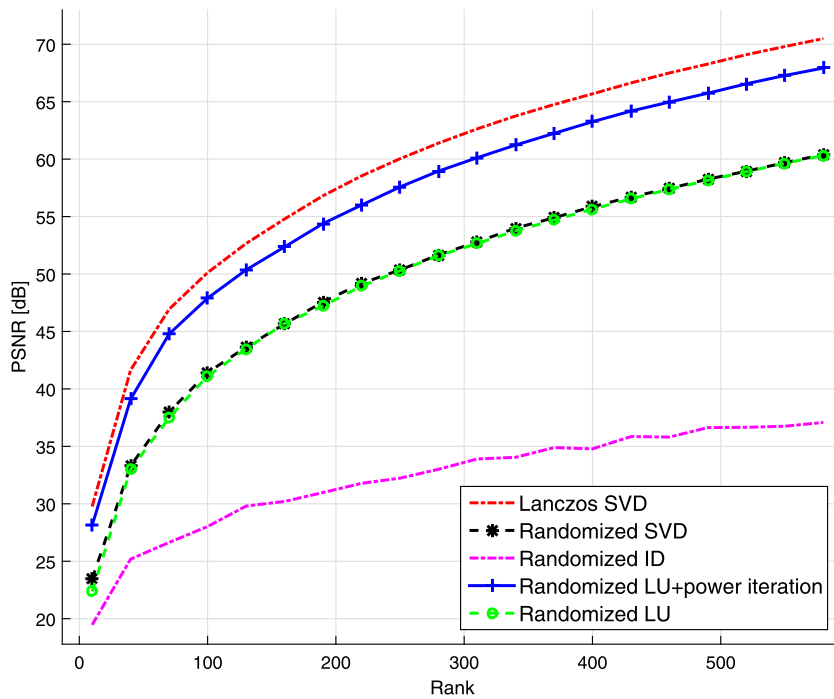


Fig. 5.9. PSNR values from image reconstruction application using randomized LU, randomized ID, randomized SVD and Lanczos SVD algorithms.

and almost identical to the randomized SVD. In addition, the PSNR values are close to the result achieved by the application of the Lanczos SVD which is the best possible rank k approximation. The execution time of each algorithm is shown in Fig. 5.10. All the computations were done in double precision. Here, the randomized LU is faster than all the other compared methods making it applicable for real time applications.

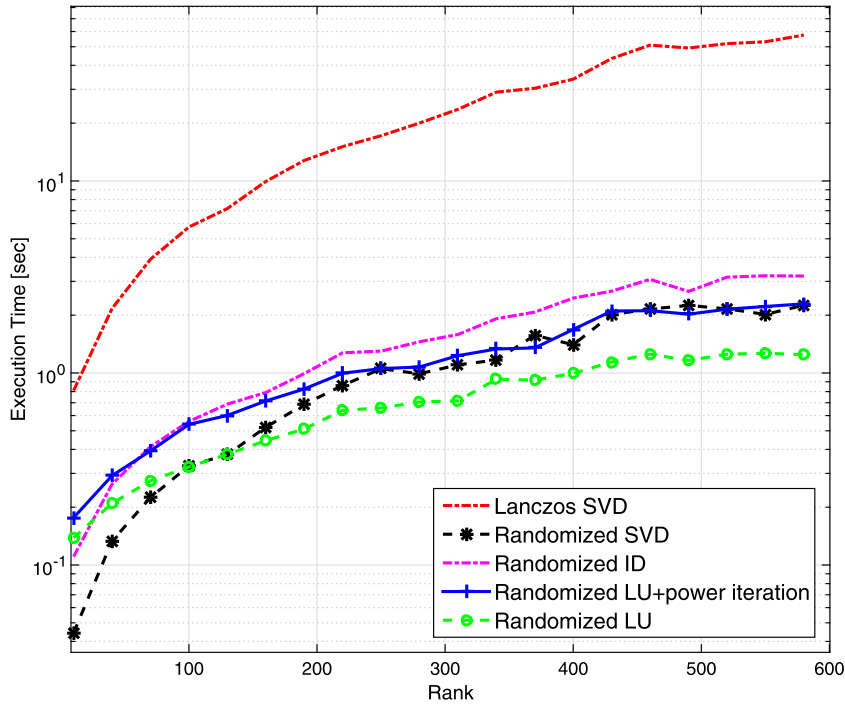


Fig. 5.10. The execution time of the randomized LU, randomized ID, randomized SVD and Lanczos SVD algorithms.




Fig. 5.11. Running time of the fast randomized LU and the randomized LU algorithms.

5.4. Fast randomized LU

In order to compare the decomposition running time for Algorithms 4.1 and 4.4, we apply these algorithms to different matrix sizes.

The y-axis in Fig. 5.11 is the time (in seconds) for decomposing an $n \times n$ matrix with $l = 3 \log_2^2 n$ where n is the x-axis.




Fig. 5.12. The normalized error (Eq. (5.1)) from the fast randomized LU and the randomized LU algorithms.




Fig. 5.13. The normalized error from the fast randomized LU and the randomized LU algorithms for slowly decaying singular values.

In addition, we see in Fig. 5.12 that the error from Algorithm 4.4 is larger than the error that Algorithm 4.1 generates. Both errors decrease at the same rate. Fig. 5.12, like Fig. 5.1, shows the relative error (Eq. (5.1)) for a randomly chosen matrix of size 3000×3000 with exponentially decaying singular values where $l = k + 3$ for different k values.

The experiment from section 5.1 was repeated, with a slowly decaying singular values. The decay of the singular values is the same as was used for the power iterations comparison $\sigma_k = \frac{100}{(9+k)^2}$. These results appear in Fig. 5.13.

The error of the fast randomized LU is larger than the error of the randomized LU because the space from which the projections are chosen is much smaller than the space created by the Gaussian-based random projection. The space is smaller since the projection matrix contains only a random diagonal matrix and

a random column selection matrix. This large error is also reflected in the error bounds of the algorithm (Theorem 4.12) and also in the need for a larger l compared to k (Lemma 4.13).

6. Conclusion

In this work, we presented a randomized algorithm for computing an LU rank k decomposition. Given an integer k , the algorithm finds an LU decomposition where both L and U are of rank k with negligible failure probability. Error bounds for the approximation of the input matrix were derived, and were proved to be proportional to the $(k + 1)$ th singular value. The performance of the algorithm (error and computational time) was compared to the randomized SVD, randomized ID and to the application of Lanczos SVD. We also showed that the algorithm can be parallelized since it consists mostly of matrix multiplication and pivoted LU. The results on GPU show that it is possible to reduce the computational time significantly even simply by using only the standard MATLAB libraries.

Acknowledgment

This research was partially supported by the Israel Science Foundation (Grant No. 1041/10), by the Israel Ministry of Science and Technology (Grants Nos. 3-9096, 3-10898), by US – Israel Binational Science Foundation (BSF 2012282), Blavatnik Computer Science Research Fund, Blavatnik ICRC Funds and by a Fellowship from Jyväskylä University.

References

- [1] G. Stewart, The decompositional approach to matrix computation, *Comput. Sci. Eng.* 2 (1) (2000) 50–59.
- [2] M. Elad, M. Aharon, Image denoising via sparse and redundant representations over learned dictionaries, *IEEE Trans. Image Process.* 15 (12) (2006) 3736–3745.
- [3] R. Mazumder, T. Hastie, R. Tibshirani, Spectral regularization algorithms for learning large incomplete matrices, *J. Mach. Learn. Res.* 99 (2010) 2287–2322.
- [4] Y. Koren, R. Bell, C. Volinsky, Matrix factorization techniques for recommender systems, *Computer* 42 (8) (2009) 30–37.
- [5] L. Wolf, A. Shashua, Learning over sets using kernel principal angles, *J. Mach. Learn. Res.* 4 (2003) 913–931.
- [6] R.R. Coifman, S. Lafon, Diffusion maps, *Appl. Comput. Harmon. Anal.* 21 (1) (2006) 5–30.
- [7] G. Shabat, Y. Shmueli, A. Bermanis, A. Averbuch, Accelerating particle filter using randomized multiscale and fast multipole type methods, *IEEE Trans. Pattern Anal. Mach. Intell.* 37 (7) (2015) 1396–1407.
- [8] O. Schenk, K. Gärtner, W. Fichtner, Efficient sparse LU factorization with left–right looking strategy on shared memory multiprocessors, *BIT Numer. Math.* 40 (1) (2000) 158–176.
- [9] J.W. Demmel, S.C. Eisenstat, J.R. Gilbert, X.S. Li, J.W. Liu, A supernodal approach to sparse partial pivoting, *SIAM J. Matrix Anal. Appl.* 20 (3) (1999) 720–755.
- [10] T.A. Davis, I.S. Duff, An unsymmetric-pattern multifrontal method for sparse LU factorization, *SIAM J. Matrix Anal. Appl.* 18 (1) (1997) 140–158.
- [11] G.H. Golub, C.F. Van Loan, *Matrix Computations*, vol. 4, John Hopkins University Press, 2012.
- [12] D. Kirk, nVidia CUDA software and GPU parallel computing architecture, in: *ISMM*, vol. 7, 2007, pp. 103–104.
- [13] P. Martinsson, V. Rokhlin, M. Tygert, A randomized algorithm for the decomposition of matrices, *Appl. Comput. Harmon. Anal.* 30 (1) (2011) 47–68.
- [14] N. Halko, P.-G. Martinsson, J.A. Tropp, Finding structure with randomness: probabilistic algorithms for constructing approximate matrix decompositions, *SIAM Rev.* 53 (2) (2011) 217–288.
- [15] A.E. Litvak, A. Pajor, M. Rudelson, N. Tomczak-Jaegermann, Smallest singular value of random matrices and geometry of random polytopes, *Adv. Math.* 195 (2) (2005) 491–523.
- [16] A. Litvak, O. Rivasplata, Smallest singular value of sparse random matrices, *Studia Math.* 212 (2010) 195–218.
- [17] A. Bermanis, A. Averbuch, R. Coifman, Multiscale data sampling and function extension, *Appl. Comput. Harmon. Anal.* 34 (2013) 15–29.
- [18] G. David, Anomaly detection and classification via diffusion processes in hyper-networks, Ph.D. thesis, School of Computer Science, Tel Aviv University, March 2009.
- [19] D.L. Donoho, Compressed sensing, *IEEE Trans. Inform. Theory* 52 (4) (2006) 1289–1306.
- [20] H. Avron, P. Maymounkov, S. Toledo, Blendenpik: supercharging LAPACK’s least-squares solver, *SIAM J. Sci. Comput.* 32 (3) (2010) 1217–1236.
- [21] T.F. Chan, Rank revealing QR factorizations, *Linear Algebra Appl.* 88 (1987) 67–82.
- [22] M. Gu, S.C. Eisenstat, Efficient algorithms for computing a strong rank-revealing QR factorization, *SIAM J. Sci. Comput.* 17 (4) (1996) 848–869.

- [23] C.-T. Pan, On the existence and computation of rank-revealing LU factorizations, *Linear Algebra Appl.* 316 (1) (2000) 199–222.
- [24] L. Miranian, M. Gu, Strong rank revealing LU factorizations, *Linear Algebra Appl.* 367 (2003) 1–16.
- [25] H. Cheng, Z. Gimbutas, P. Martinsson, V. Rokhlin, On the compression of low rank matrices, *SIAM J. Sci. Comput.* 26 (4) (2005) 1389–1404.
- [26] P. Drineas, M.W. Mahoney, S. Muthukrishnan, Relative-error CUR matrix decompositions, *SIAM J. Matrix Anal. Appl.* 30 (2) (2008) 844–881.
- [27] V. Rokhlin, M. Tygert, A fast randomized algorithm for overdetermined linear least-squares regression, *Proc. Nat. Acad. Sci.* 105 (36) (2008) 13212–13217.
- [28] K.L. Clarkson, D.P. Woodruff, Low rank approximation and regression in input sparsity time, in: *Proceedings of the 45th Annual ACM Symposium on Symposium on Theory of Computing*, ACM, 2013, pp. 81–90.
- [29] D. Achlioptas, F. McSherry, Fast computation of low-rank matrix approximations, *J. ACM* 54 (2) (2007) 9.
- [30] K.L. Clarkson, D.P. Woodruff, Numerical linear algebra in the streaming model, in: *Proceedings of the 41st Annual ACM Symposium on Theory of Computing*, ACM, 2009, pp. 205–214.
- [31] A. Magen, A. Zouzias, Low rank matrix-valued Chernoff bounds and approximate matrix multiplication, in: *Proceedings of the Twenty-Second Annual ACM–SIAM Symposium on Discrete Algorithms*, SIAM, 2011, pp. 1422–1436.
- [32] A. Frieze, R. Kannan, S. Vempala, Fast Monte-Carlo algorithms for finding low-rank approximations, *J. ACM* 51 (6) (2004) 1025–1041.
- [33] P. Drineas, R. Kannan, M.W. Mahoney, Fast Monte Carlo algorithms for matrices II: computing a low-rank approximation to a matrix, *SIAM J. Comput.* 36 (1) (2006) 158–183.
- [34] C. Boutsidis, A. Gittens, Improved matrix algorithms via the subsampled randomized Hadamard transform, *SIAM J. Matrix Anal. Appl.* 34 (3) (2013) 1301–1340.
- [35] R. Bhatia, *Matrix Analysis*, vol. 169, Springer, 1997.
- [36] H.H. Goldstine, J. Von Neumann, Numerical inverting of matrices of high order II, *Proc. Amer. Math. Soc.* 2 (2) (1951) 188–202.
- [37] M. Rudelson, R. Vershynin, Smallest singular value of a random rectangular matrix, *Comm. Pure Appl. Math.* 62 (12) (2009) 1707–1739.
- [38] N. Ailon, B. Chazelle, The fast Johnson–Lindenstrauss transform and approximate nearest neighbors, *SIAM J. Comput.* 39 (1) (2009) 302–322.
- [39] F. Woolfe, E. Liberty, V. Rokhlin, M. Tygert, A fast randomized algorithm for the approximation of matrices, *Appl. Comput. Harmon. Anal.* 25 (3) (2008) 335–366.
- [40] L.N. Trefethen, R.S. Schreiber, Average-case stability of Gaussian elimination, *SIAM J. Matrix Anal. Appl.* 11 (3) (1990) 335–360.
- [41] G. Stewart, The triangular matrices of Gaussian elimination and related decompositions, Tech. rep. TR-3533, Department of Computer Science and Institute for Advanced Computer Studies, University of Maryland, College Park, MD, 1995.
- [42] Z. Chen, J.J. Dongarra, Condition numbers of Gaussian random matrices, *SIAM J. Matrix Anal. Appl.* 27 (3) (2005) 603–620.
- [43] R. Witten, E. Candes, Randomized algorithms for low-rank matrix factorizations: sharp performance bounds, *Algorithmica* (2013) 1–18.
- [44] J.A. Tropp, Improved analysis of the subsampled randomized Hadamard transform, *Adv. Adapt. Data Anal.* 3 (01–02) (2011) 115–126.
- [45] P.-G. Martinsson, A. Szlam, M. Tygert, Normalized power iterations for the computation of SVD, Manuscript, Nov.
- [46] R. Larsen, Lanczos bidiagonalization with partial reorthogonalization, Tech. rep. DAIMI PB-357, Department of Computer Science, Aarhus University, 1998.